

# CS525-04/05: Advanced Database Organization

## Notes 1: Introduction to DBMS Implementation

Yousef M. Elmehdwi

Department of Computer Science

Illinois Institute of Technology

[yelmehdwi@iit.edu](mailto:yelmehdwi@iit.edu)

August 23<sup>rd</sup> 2023

Slides: adapted from a course taught by [Hector Garcia-Molina](#), Stanford

# Core Terminology Review

- Data
  - Data refers to any piece of information that holds value and is worth keeping.
  - It's often stored in electronic form and can range from numbers and text to images and videos.
- Database
  - organized collection of interrelated data that models some aspect of the real-world.
- Query
  - operation that retrieves specific data from a database based on certain criteria or conditions.
  - queries allow users to extract relevant information.
- Relation
  - refers to the organization of data into a two-dimensional table, where rows (tuples) represent basic entities or facts of some sort, and columns (attributes) represent properties of those entities.
- Schema
  - a description of the structure of the data in a database, often called "metadata"
  - it's like a blueprint that outlines how the data is organized, what types of data are stored, and how they are related.

# Database Management System (DBMS)

- A DBMS is software that allows applications to store and analyze information in a database.
- A general-purpose DBMS is designed to allow the definition, creation, querying, update, and administration of databases.

# Advanced Database Organization?

- =Database Implementation
- =How to implement a database system
- and have fun doing it ;-)

# What do you want from a DBMS?

- Keep data around (persistent)
- Answer questions (queries) about data
- Update data

# Isn't Implementing a Database System Simple?

- Relation  $\Rightarrow$  Statements  $\Rightarrow$  Results

# Introduction the MEGATRON 3000 Database Management System

- “Imaginary” database System
- The latest from MEGATRON Labs
- Incorporates latest relational technology
- UNIX compatible
- Lightweight & cheap!

# MEGATRON 3000 Implementation Details

- MEGATRON 3000 uses the file system to store its relations
- Relations stored in files (ASCII)
- Use a separate file per entity/relation.
- The application has to parse the files each time they want to read/update records.
  - e.g., relation **Students**(*name, id, dept*) is in `/usr/db/Students`
  - The file **Students** has one line for each tuple.
  - Values of components of a tuple are stored as a character string, separated by special marker character #

|        |   |     |   |    |
|--------|---|-----|---|----|
| Smith  | # | 123 | # | CS |
| Jonson | # | 522 | # | EE |
|        |   | :   |   |    |

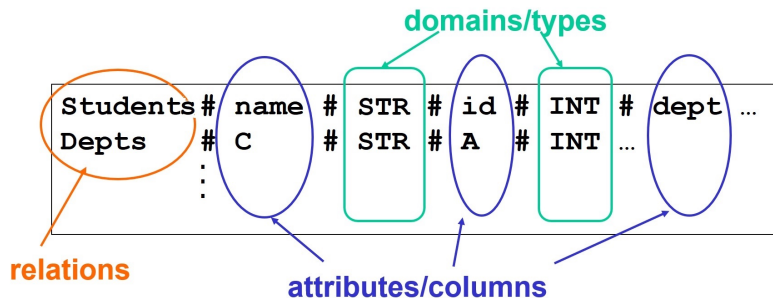


# MEGATRON 3000 Implementation Details

- The database schema is stored in a special file
- Schema file (ASCII) in `/usr/db/schema`
  - For each relation, the file `schema` has a line beginning with that relation name, in which attribute names alternate with types.
  - The character `#` separates elements of these lines.

|          |   |      |   |     |   |    |   |     |     |          |
|----------|---|------|---|-----|---|----|---|-----|-----|----------|
| Students | # | name | # | STR | # | id | # | INT | #   | dept ... |
| Depts    | # | C    | # | STR | # | A  | # | INT | ... |          |
|          |   |      |   | :   |   |    |   |     |     |          |

# MEGATRON 3000 Implementation Details



# MEGATRON 3000 Sample Sessions

```
% MEGATRON3000
    Welcome to MEGATRON 3000!
&
:
& quit
%
```

- We are now talking to the MEGATRON 3000 user interface, to which we can type SQL queries in response to the Megatron prompt (&).

# MEGATRON 3000 Sample Sessions

```
& select * from Students #
```

Relation Students

name      id      dept

Smith      123      CS

Johnson 522      EE

...

&

columns/attributes

rows/tuples

- A # ends a query

# MEGATRON 3000 Sample Sessions

- Execute a query and send the result to printer

```
& select *  
  from Students | LPR #  
&
```

- Result sent to LPR (printer).

# MEGATRON 3000 Sample Sessions

- Execute a query and store the result in a new file

```
& select *  
  from Students  
 where id < 100 | LowId #  
&
```

- New relation LowId created.

# How MEGATRON 3000 Executes Queries

- To execute

```
SELECT * FROM R WHERE <condition>
```

- 1 Read schema to get attributes of R
- 2 Check validity of condition
- 3 Display attributes of R as the header
- 4 Read file R; for each line:
  - a Check condition
  - b If TRUE, display the line as tuple

# MEGATRON 3000 Query Execution

- To execute

```
SELECT * FROM R WHERE <condition> | T
```

- 1 Process select as before but omit Step 3
- 2 Write results to new file T
- 3 Append new line to dictionary `usr/db/schema`



# MEGATRON 3000 Query Execution

- Consider a more complicated query, one involving a join of two relations R, S
- To execute

```
SELECT A,B FROM R,S WHERE <condition>
```

- 1 Read schema to get R,S attributes
- 2 Read R file, for each line *r*:
  - 3 Read S file, for each line *s*:
    - 1 Create join tuple *r* & *s*
    - 2 Check condition
    - 3 If TRUE, Display *r,s*[A,B]

# What's wrong with MEGATRON 3000 DBMS?

- DBMS is not implemented like our *imaginary* MEGATRON 3000
- Described implementation is inadequate for applications involving significant amount of data or multiple users of data
- Partial list of problems follows

# What's wrong with MEGATRON 3000 DBMS?

- Tuple layout on disk is inadequate with no flexibility when the database is modified
- e.g., change String from *CS* to *CSDept* in one **Students** tuple, we have to rewrite the entire file
  - ASCII storage is expensive
  - Deletions are expensive

# What's wrong with MEGATRON 3000 DBMS?

- Search expensive; no indexes
  - e.g., cannot find tuple with given key quickly
  - Always have to read full relation

# What's wrong with MEGATRON 3000 DBMS?

- Brute force query processing
- e.g.,

```
SELECT * FROM R,S WHERE R.A = S.A and S.B > 1000
```

- Much better if use index to select tuples that satisfy condition (Do select using  $S.B > 1000$  first)
- More efficient join (sort both relations on  $A$  and merge)

# What's wrong with MEGATRON 3000 DBMS?

- No buffer manager
  - There is no way for useful data to be buffered in main memory; all data comes off the disk, all the time
  - e.g., need caching.

# What's wrong with MEGATRON 3000 DBMS?

- No concurrency control
  - Several users can modify a file at the same time with unpredictable results.

# What's wrong with MEGATRON 3000 DBMS?

- No reliability
- e.g., in case of error/crash, say, power failure or leave operations half done
  - Can lose data



# What's wrong with MEGATRON 3000 DBMS?

- No security
- e.g., file system security is coarse
  - Unable to restrict access, say, to some fields of a relation and not others

# What's wrong with MEGATRON 3000 DBMS?

- No application program interface (API)
  - e.g., how can a payroll program get at the data?

# What's wrong with MEGATRON 3000 DBMS?

- Cannot interact with other DBMSs.

# What's wrong with MEGATRON 3000 DBMS?

- No GUI

# This Course

- Introduce students to better way of building a database management systems.

# Reading assignment

- Refresh your memory about basics of the relational model and SQL
  - from your earlier course notes
  - from some textbook
  - <http://cs.iit.edu/~cs425/schedule.html>



- Course Blackboard: Assignments\Reading subfolder
  - Chapter 1: “Introduction to DBMS Implementation”

Notes 2: Hardware