

CS525: Advanced Database Organization

Notes 3: Database Storage Part II

Yousef M. Elmehdwi

Department of Computer Science

Illinois Institute of Technology

yelmehdwi@iit.edu

September 6th 2023

Slides: adapted from courses taught by [Andy Pavlo](#), Carnegie Mellon University, [Hector Garcia-Molina](#), Stanford, & [Shun Yan Cheung](#), Emory University

- Database Systems: The Complete Book, 2nd Edition,
 - *Chapter 2: Data Storage*
- Database System Concepts (6th/7th Edition)
 - *Chapter 10 (6th)/ Chapter 13 (7th)*

Database Storage

- **Problem#1:** How the DBMS represents the database in files on disk
 - i.e., how to lay out data on disk.
- **Problem#2:** How the DBMS manages its memory and move data back-and-forth from disk.

Today's Agenda

- Data Representation

- *How the system stores the actual binary data for individual attributes (columns) within the database.*

- System Catalogs

- *Internal metadata is maintained by the database to understand both the data that is actually stored and how to interpret the bytes within the tuples*

- Storage Models

- *How data is organized and stored within the database system.*

- Modification of Tuples


Tuple Storage

- A **tuple** is essentially a sequence of bytes (byte arrays).
- It is up to the DBMS to know how to interpret those bytes to derive the values for attributes.
- The **DBMS's catalogs** contain the schema information about tables that the system uses to figure out the tuple's layout.

What are the data items we want to store?

- a salary
- a name
- a date
- a picture

⇒ What we have available: Bytes



Data Representation

- How a DBMS stores the bytes for a value?
 - *How the DBMS stores the binary data for different types of values or attributes?*
- There are five high level data types that can be stored in tuples:
 - integers,
 - variable precision numbers,
 - fixed point precision numbers,
 - variable length values, and
 - dates/times.

IEEE-754 Standard¹

- This is a specific standard that defines the binary representation of floating-point numbers (like float and double) and their arithmetic operations in computing.
- Ensures consistency in how floating-point numbers are stored and manipulated across different computer architectures and programming languages.
- It specifies the format for representing real numbers, including the sign bit, exponent, and mantissa (fractional part).
- Example: 32 bit in Standard IEEE 754:



To be represented in this format, a number should be in the following normalized form.
 $(+ \text{ or } -) 1.(\text{mantissa}) \times 2^{(\text{exponent})}$

¹ https://en.wikipedia.org/wiki/IEEE_754

Data Representation

How a DBMS stores the bytes for a value

- **INTEGER/BIGINT/SMALLINT/TINYINT**
 - C/C++ Representation
 - All integers are stored in their “native” C/C++ types¹ within the database.
- **FLOAT/REAL** vs. **NUMERIC/DECIMAL**
 - IEEE-754 Standard/Fixed-point Decimals
- **VARCHAR/VARBINARY/TEXT/BLOB**
 - Header with length, followed by data bytes.
- **TIME/DATE/TIMESTAMP**
 - 32/64-bit integer of (micro)seconds since Unix epoch

¹refer to the fundamental data types that are supported directly by the C/C++ programming languages without any additional libraries or custom data types. These native data types are typically used for storing and manipulating data efficiently in these languages.

Data Representation: Integers

- C/C++ Representation
 - Most DBMSs store integers using their “native” C/C++ types as specified by the [IEEE-754 standard](#).
- These values are fixed length.
- Examples: [INTEGER](#)/[BIGINT](#)/[SMALLINT](#)/[TINYINT](#)

Variable Precision Numbers

- These are inexact¹, variable-precision numeric types that uses the “native” C/C++ types.
- Store directly as specified by [IEEE-754 standard](#).
- These values are also fixed length.
- Typically faster than arbitrary precision numbers because the CPU can execute instructions on them directly.
 - Example: `FLOAT`, `REAL/DOUBLE`
- but can have **rounding errors**² when performing computations due to the fact that some numbers cannot be represented precisely in binary floating-point format.
- As a result, calculations may yield slightly inaccurate results.
 - To avoid this issue, we use Fixed-Point Precision Numbers.

¹ *Inexact means that some values cannot be converted exactly to the internal format and are stored as approximations, so that storing and retrieving a value might show slight discrepancies.*

² *rounding error, is the difference between the result produced by a given algorithm using exact arithmetic and the result produced by the same algorithm using finite-precision, rounded arithmetic.*

Variable Precision Numbers

- Rounding Example

```
#include <stdio.h>

int main(int argc, char* argv[]) {
    float x = 0.1;
    float y = 0.2;
    printf("x+y = %f\n", x+y);
    printf("0.3 = %f\n", 0.3);
}
```

Output

```
x+y = 0.300000
0.3 = 0.300000
```

Variable Precision Numbers

- Rounding Example

```
#include <stdio.h>

int main(int argc, char* argv[]) {
    float x = 0.1;
    float y = 0.2;
    printf("x+y = %.20f\n", x+y);
    printf("0.3 = %.20f\n", 0.3);
}
```

Output

```
x+y = 0.30000001192092895508
0.3 = 0.29999999999999998890
```

Variable Precision Numbers

- Rounding Example

```
public class RoundingError {  
    public static void main(String[] args) {  
        float x = 1.0f;  
        for (int i = 0; i < 10; i++) {  
            x -= 0.1f;  
        }  
        System.out.printf("Result w precision: %.1f\n", x);  
        System.out.printf("Result w precision 10: %.10f\n", x);  
    }  
}
```

Output

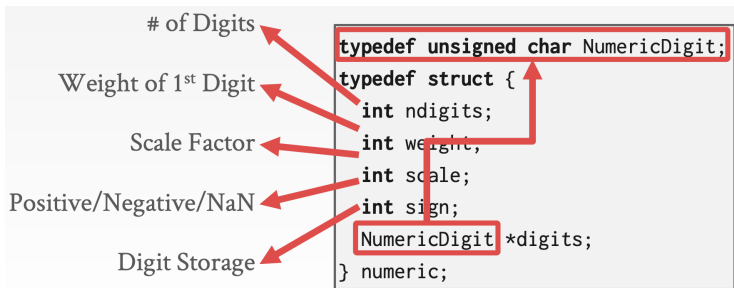
Result w precision 1 : -0.0

Result w precision 10: -0.0000000745

Data Representation: Fixed Point Precision Numbers

- Numeric data types with arbitrary precision and scale.
- Used when round errors are unacceptable.
 - Example: **NUMERIC**, **DECIMAL**
- Typically stored in an exact, variable-length binary representation with additional meta-data that specifies the length of the data and the position of the decimal point.
 - Like a **VARCHAR** but not stored as a string
- but the DBMS pays a performance penalty to get this accuracy.
 - *Calculations involving fixed-point precision numbers may be slower compared to operations with native numeric types like **FLOAT** or **DOUBLE**, which use hardware-based floating-point arithmetic.*

PostgreSQL: NUMERIC




```

/*
 * add_var() -
 *
 * Full version of add functionality on variable level (handling signs).
 * result might point to one of the operands too without danger.
 */
static void
add_var(const NumericVar *var1, const NumericVar *var2, NumericVar *result)
{
    /*
     * Decide on the signs of the two variables what to do
     */
    if (var1->sign == NUMERIC_POS)
    {
        if (var2->sign == NUMERIC_POS)
        {
            /*
             * Both are positive result = +(ABS(var1) + ABS(var2))
             */
            add_abs(var1, var2, result);
            result->sign = NUMERIC_POS;
        }
        else
        {
            /*
             * var1 is positive, var2 is negative Must compare absolute values
             */
            switch (cmp_abs(var1, var2))
            {
                case 0:
                    /*
                     * ABS(var1) == ABS(var2)
                     * result = ZERO
                     */
                    zero_var(result);
                    result->dscale = Max(var1->dscale, var2->dscale);
                    break;

                case 1:
                    /*
                     * ABS(var1) > ABS(var2)
                     * result = +(ABS(var1) - ABS(var2))
                     */
                    sub_abs(var1, var2, result);
                    result->sign = NUMERIC_POS;
                    break;

                case -1:
                    /*
                     * ABS(var1) < ABS(var2)
                     * result = -(ABS(var2) - ABS(var1))
                     */
                    sub_abs(var2, var1, result);
                    result->sign = NUMERIC_NEG;
                    break;
            }
        }
    }
    else
    {
        if (var2->sign == NUMERIC_POS)
        {
            /*
             * var1 is negative, var2 is positive
             * Must compare absolute values
             */
            switch (cmp_abs(var1, var2))
            {
                case 0:
                    /*
                     * ABS(var1) == ABS(var2)
                     * result = ZERO
                     */
                    zero_var(result);
                    result->dscale = Max(var1->dscale, var2->dscale);
                    break;
            }
        }
    }
}

```

Data Representation: Variable Length Data

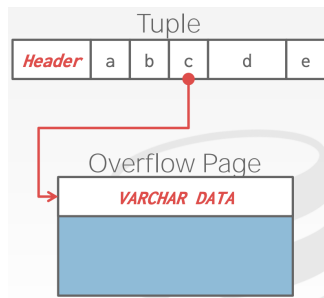
- These represent data types of arbitrary length.
 - An array of bytes of arbitrary length.
- Has a header that keeps track of the length of the string to make it easy to jump to the next value. It may also contain a checksum for the data.
- Example: `VARCHAR`, `VARBINARY`, `TEXT`, `BLOB`.

Large Values

- Most DBMSs don't allow a tuple to exceed the size of a single page.
- Handling tuples that exceed the size of a single page in a DBMS is a common challenge.
- DBMSs typically have strategies to deal with this situation to ensure data integrity and efficient storage.
- Two common approaches to handle such cases are:
 - overflow page
 - external storage

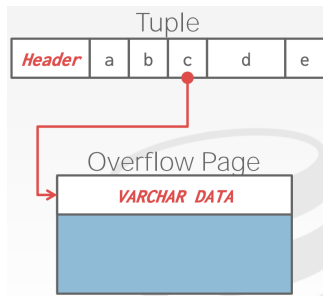
Large Values: Overflow Page

- To store values that are larger than a page, the DBMS uses separate **overflow** storage pages and have the tuple contain a reference to that page.
 - The main part of the tuple, which can fit within a single page, is stored in the primary page, while the overflowed part is stored in one or more additional overflow pages.
 - Overflow pages are linked to the primary page, forming a chain of pages that together represent the complete tuple.
 - When querying the data, the DBMS follows these chains of pages to reconstruct the complete tuple.
- These overflow pages can contain pointers to additional overflow pages until all the data can be stored.




Large Values: Overflow Page

- To store values that are larger than a page, the DBMS uses separate **overflow** storage pages and have the tuple contain a reference to that page.
- Different DBMSs have different name/specification/requirements when they do that:
 - Postgres: TOAST (The Oversized-Attribute Storage Technique) (>2KB)
 - MySQL: Overflow (> $\frac{1}{2}$ size of page)
 - SQL Server: Overflow (> size of page)
- These overflow pages can contain pointers to additional overflow pages until all the data can be stored.



External Value Storage

- Some systems allow to store large data values, such as files or binary objects, in an **external file** rather than directly within the database, and then the tuple will contain a pointer to that file.
- Example:
 - if the database is storing photo information, the DBMS can store the photos in the external files rather than having them take up large amounts of space in the DBMS.
- Treated as a **BLOB** type
 - Oracle: **BFILE** data type
 - Contains a locator pointing to a large binary file stored outside the database.
 - Microsoft: **FILESTREAM** data type
- The DBMS **cannot** manipulate the contents of an external file.
- Reading: A paper explains the trade-offs between these two options:
 - To BLOB or Not To BLOB: Large Object Storage in a Database or a Filesystem



Data Representation: Dates and Times

- Varies widely across different database systems?
- However, a common approach is to represent dates and times as the number of (micro/milli)seconds since the Unix epoch¹.
- Example: **TIME**, **DATE**, **TIMESTAMP**.

¹ *The Unix epoch is a reference point in time, representing January 1, 1970, at 00:00:00 UTC (Coordinated Universal Time). It is widely used as a starting point for measuring time intervals.*

System Catalog

- In order for the DBMS to be able to decipher the contents of tuples, it maintains an internal catalog to tell it meta-data about the databases
 - A DBMS stores meta-data about databases in its internal catalogs.
- The meta-data will contain what tables and columns the databases have along with their types and the orderings of the values.
 - Tables, columns, indexes, views
 - Users, permissions
 - Internal statistics
- Almost every DBMS stores their a database's catalog in itself in the format that they use for their tables
 - They use special code to **bootstrap**¹ these catalog tables (wrap low-level access methods to access the catalog)

¹ *Bootstrapping is the process of initializing a DBMS's catalog during system setup or database creation. During this phase, the DBMS uses low-level access methods or internal mechanisms to create the catalog tables and populate them with initial data*

- You can query the DBMS's internal `INFORMATION_SCHEMA` catalog to get info about the database.
 - ANSI standard set of read-only views that provide info about all of the tables, views, columns, and procedures in a database.
- DBMSs also have non-standard shortcuts to retrieve this information.

Accessing Table Schema

- List all of the tables in the current database:

```
-- SQL-92
SELECT *
FROM INFORMATION_SCHEMA.TABLES
WHERE table_catalog = '<db name>';
```

```
\d;           -- Postgres
SHOW TABLES; -- MySQL
.tables;      -- SQLite
```

Accessing Table Schema

- List all of the columns in the *student* table:

```
-- SQL-92
SELECT *
FROM INFORMATION_SCHEMA.TABLES
WHERE table_name = 'student'
```

```
\d student;          -- Postgres
DESCRIBE student;    -- MySQL
.schema student;     -- SQLite
```

Today's Agenda

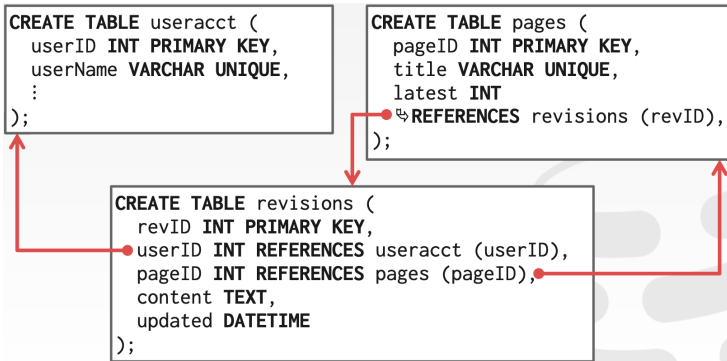
- Data Representation
- System Catalogs
- Storage Models
 - Ways to store tuples in pages

Observation

- The relational model does **not** specify that we have to store all of a tuple's attributes together in a single page.
- This may not actually be the best layout for some workloads
- There are many different workloads for database systems.
- By **workload**¹, we are referring to the general nature of requests a system will have to handle.
- Different workloads have different requirements for data storage and access patterns.

¹ *refers to the types of queries, transactions, and operations the system is expected to handle.*

Wikipedia Example



- On-line Transaction Processing:
 - Simple queries that read/update a small amount of data that is related to a single entity in the database.
- This is usually the kind of application that people build first.

```
SELECT P.*, R.*  
FROM pages AS P  
      INNER JOIN revisions AS R  
      ON P.latest = R.revID  
WHERE P.pageID = ?
```

```
UPDATE useracct  
SET lastLogin = NOW(),  
    hostname = ?  
WHERE userID = ?
```

```
INSERT INTO revisions  
VALUES (?,?...,?)
```

OLTP: On-line Transaction Processing

- Fast, short running operations
- Simple queries that operate on single entity at a time
- Typically handle more writes than reads
- Repetitive operations
- Usually the kind of application that people build first
- Example
 - User invocations of Amazon (Amazon storefront).
 - Users can add things to their cart,
 - they can make purchases,
 - but the actions only affect their accounts.

- On-line Analytical Processing:
 - Complex queries that read large portions of the database spanning multiple entities.
- You execute these workloads on the data you have collected from your OLTP application(s).

```
SELECT COUNT(U.lastLogin),  
       EXTRACT(month FROM  
              U.lastLogin) AS month  
FROM useracct AS U  
WHERE U.hostname LIKE '%.gov'  
GROUP BY  
       EXTRACT(month FROM U.  
              lastLogin)
```

OLAP: On-line Analytical Processing

- Long running, more complex queries
- Reads large portions of the database
- Analyzing and deriving new data from existing data collected on the OLTP side
- Example
 - Amazon computing the five most bought items over a one month period for these geographical locations.

HTAP: Hybrid Transaction + Analytical Processing

- A new type of workload which has become popular recently is HTAP, which is like a combination which tries to do OLTP and OLAP together on the same database.
- Watch [HTAP Databases: What is New and What is Next - SIGMOD22-HTAP-Tutorial- June 2022](#)

Data Storage Model

- There are different ways to store tuples in pages.
- The DBMS can store tuples in different ways that are better for either OLTP or OLAP workloads
- We have been assuming the *n-ary storage model* (aka “row storage”) so far this semester.

N-ARY Storage Model (NSM)

- The DBMS stores all attributes for a single tuple contiguously in a single page.
- Ideal for OLTP workloads where requests are insert-heavy and transactions tend to operate only on an individual entity
 - it takes only one fetch to be able to get all of the attributes for a single tuple.

N-ARY Storage Model (NSM)

- The DBMS stores all attributes for a single tuple contiguously in a single page.

<i>Header</i>	userID	userName	userPass	hostname	lastLogin	Tuple #1
<i>Header</i>	userID	userName	userPass	hostname	lastLogin	
<i>Header</i>	userID	userName	userPass	hostname	lastLogin	
<i>Header</i>	-	-	-	-	-	


N-ARY Storage Model (NSM)

- The DBMS stores all attributes for a single tuple contiguously in a single page.

	<i>Header</i>	userID	userName	userPass	hostname	lastLogin	Tuple #1
Tuple #2	<i>Header</i>	userID	userName	userPass	hostname	lastLogin	
	<i>Header</i>	userID	userName	userPass	hostname	lastLogin	Tuple #3
Tuple #4	<i>Header</i>	-	-	-	-	-	

N-ARY Storage Model (NSM)

- The DBMS stores all attributes for a single tuple contiguously in a single page.



N-ARY Storage Model (NSM)

```
SELECT * FROM useracct  
WHERE userName = ? AND userPass = ?
```

```
INSERT INTO useracct  
VALUES (?, ?, ..., ?)
```

N-ARY Storage Model (NSM)

```
SELECT * FROM useracct  
WHERE userName = ?  
AND userPass = ?
```




NSM Disk Page

Header	userID	userName	userPass	hostname	lastLogin
Header	userID	userName	userPass	hostname	lastLogin
Header	userID	userName	userPass	hostname	lastLogin
Header	-	-	-	-	-

N-ARY Storage Model (NSM)

```
SELECT * FROM useracct  
WHERE userName = ?  
AND userPass = ?
```

```
INSERT INTO useracct  
VALUES (?, ?, ...?)
```



NSM Disk Page

Header	userID	userName	userPass	hostname	lastLogin
Header	userID	userName	userPass	hostname	lastLogin
Header	userID	userName	userPass	hostname	lastLogin
Header	userID	userName	userPass	hostname	lastLogin

N-ARY Storage Model (NSM)

```
SELECT COUNT(U.lastLogin),  
       EXTRACT(month FROM U.lastLogin) AS month  
FROM   useracct AS U  
WHERE  U.hostname LIKE '%.gov'  
GROUP BY EXTRACT(month FROM U.lastLogin)
```

N-ARY Storage Model (NSM)

```
SELECT COUNT(U.lastLogin),  
       EXTRACT(month FROM U.lastLogin) AS month  
FROM useracct AS U  
WHERE U.hostname LIKE '%.gov'  
GROUP BY EXTRACT(month FROM U.lastLogin)
```



NSM Disk Page

<i>Header</i>	userID	userName	userPass	hostname	lastLogin
<i>Header</i>	userID	userName	userPass	hostname	lastLogin
<i>Header</i>	userID	userName	userPass	hostname	lastLogin
<i>Header</i>	userID	userName	userPass	hostname	lastLogin

N-ARY Storage Model (NSM)

```
SELECT COUNT(U.lastLogin),  
       EXTRACT(month FROM U.lastLogin) AS month  
FROM useracct AS U  
WHERE U.hostname LIKE '%.gov'  
GROUP BY EXTRACT(month FROM U.lastLogin)
```



NSM Disk Page

Header	userID	userName	userPass	hostname	lastLogin
Header	userID	userName	userPass	hostname	lastLogin
Header	userID	userName	userPass	hostname	lastLogin
Header	userID	userName	userPass	hostname	lastLogin

N-ARY Storage Model (NSM)

```
SELECT COUNT(U.lastLogin),  
       EXTRACT(month FROM U.lastLogin) AS month  
FROM   useracct AS U  
WHERE  U.hostname LIKE '%.gov'  
GROUP BY EXTRACT(month FROM U.lastLogin)
```




NSM Disk Page

Header	userID	userName	userPass	hostname	lastLogin
Header	userID	userName	userPass	hostname	lastLogin
Header	userID	userName	userPass	hostname	lastLogin
Header	userID	userName	userPass	hostname	lastLogin

N-ARY Storage Model (NSM)

```
SELECT COUNT(U.lastLogin),  
       EXTRACT(month FROM U.lastLogin) AS month  
FROM useracct AS U  
WHERE U.hostname LIKE '%.gov'  
GROUP BY EXTRACT(month FROM U.lastLogin)
```



Useless Data

N-ARY Storage Model (NSM)

- Advantages

- Fast inserts, updates, and deletes.
- Good for queries that need the entire tuple.

- Disadvantages

- Not good for scanning large portions of the table and/or a subset of the attributes.
- This is because it pollutes the buffer pool by fetching data that is not needed for processing the query.

Decomposition Storage Model (DSM)

- The DBMS stores the values of a single attribute (column) for all tuples contiguously in a block of data.
 - *Vertically partition a database into a collection of individual columns that are stored separately*
 - Also known as a “column store”.
- Ideal for OLAP workloads where read-only queries perform large scans over a subset of the table’s attributes.


Decomposition Storage Model (DSM)

- The DBMS stores the values of a single attribute for all tuples contiguously in a page.
 - Also known as a “column store”.

<i>Header</i>	userID	userName	userPass	hostname	lastLogin
<i>Header</i>	userID	userName	userPass	hostname	lastLogin
<i>Header</i>	userID	userName	userPass	hostname	lastLogin
<i>Header</i>	userID	userName	userPass	hostname	lastLogin

Decomposition Storage Model (DSM)

- The DBMS stores the values of a single attribute for all tuples contiguously in a page.
 - Also known as a “column store”.




Decomposition Storage Model (DSM)

```
SELECT COUNT(U.lastLogin),  
        EXTRACT(month FROM U.lastLogin) AS month  
FROM    useracct AS U  
WHERE   U.hostname LIKE '%.gov'  
GROUP BY EXTRACT(month FROM U.lastLogin)
```

Decomposition Storage Model (DSM)

```
SELECT COUNT(U.lastLogin),  
       EXTRACT(month FROM U.lastLogin) AS month  
FROM   useracct AS U  
WHERE  U.hostname LIKE '%.gov'  
GROUP BY EXTRACT(month FROM U.lastLogin)
```



- If there is a match on one page, how can we figure out a match on another page?*

Decomposition Storage Model: Tuple Identification

- To put the tuples back together when we are using a column store, we can use:
 - Choice #1: Fixed-length Offsets (most commonly used approach)
 - Choice #2: Embedded Tuple Ids (less common approach)
- *When decomposing a relational database into a column store model, it's essential to have a mechanism to identify and reconstruct the original tuples when needed.*

Choice #1: Fixed-length Offsets

- Each column of the table is stored separately, and for each column, a fixed-length offset is used to locate the position of each tuple within that column.
 - Assuming the attributes are all fixed-length, the DBMS can compute the offset of the attribute for each tuple.
- When the system wants the attribute for a specific tuple, it knows how to jump to that spot in the file from the offset.
- To accommodate the variable-length fields, the system can either pad fields so that they are all the same length or use a dictionary that takes a fixed-size integer and maps the integer to the value.

Offsets

	A	B	C	D
0				
1				
2				
3				

Embedded Ids

	A	B	C	D
0				
1				
2				
3				

Choice #2: Embedded Tuple Ids

- A less common approach
- Each column is still stored separately, but instead of using fixed-length offsets, each column contains embedded tuple IDs or pointers that indicate the position or identity of the tuple to which each value belongs.
- These embedded IDs link the values across columns, allowing the DBMS to reconstruct tuples by following the tuple IDs across columns.
- Note that this method has a large storage overhead because it needs to store a tuple id for every attribute entry.

Offsets

	A	B	C	D
0				
1				
2				
3				

Embedded Ids

	A	B	C	D
0				
1				
2				
3				

Decomposition Storage Model (DSM)

- Advantages

- Reduces the amount wasted I/O because the DBMS only reads the data that it needs.
- Enable better query processing and data compression.
 - because all of the values for the same attribute are stored contiguously

- Disadvantages

- Slow for point queries, inserts, updates, and deletes because of tuple splitting/stitching.

Modification of Tuples


How to handle the following operations on the tuple level?

- ➊ Insertion
- ➋ Deletion
- ➌ Update

1) Insertion

- **Easy case** Tuples with Fixed Length and not in Sequence (Unordered)
 - Insert new tuple at end of file
 - or, in deleted slot if one is available.
 - Deleted slots can be used to reuse space efficiently.
- **A little harder:** Tuples with Variable Size
 - If records are variable size, insertion becomes more complex.
 - Reusing space may not always be possible, leading to fragmentation where empty spaces are scattered throughout the storage.
 - This can impact storage efficiency over time.
- **A Difficult case:** Tuples in Sequence (Ordered)
 - Inserting tuples into a sequence can be challenging
 - Find position and slide following tuples to make room
 - If tuples are sequenced by linking, insert overflow blocks (there isn't enough space in the current block or node)
- *Overflow blocks are essentially additional blocks or nodes that are linked to the main sequence. They hold excess data when the main sequence's blocks are full.*

2) Deletion



Options

- (a) Deleted and immediately reclaim space by shifting other tuples or removing overflows
- (b) Mark deleted and list as free for re-use
 - May need chain of deleted tuples (for re-use)
 - Need a way to mark

Trade-offs


- How expensive is immediate reclaim?
 - How expensive is to move valid tuple to free space for immediate reclaim?
 - How much space is wasted?
 - e.g., deleted tuples, delete fields, ...
- *Space Efficiency vs. Performance: Immediate reclamation is more space-efficient but can be performance-intensive. Marking as deleted is less performance-intensive but can lead to space wastage.*

Concern with deletions

A caveat when using physical addresses to reference a block/record

Example

- Record Y can be referenced by other tuples (e.g., tuples X1 & X2)




Concern with deletions

A caveat when using physical addresses to reference a block/record

Example

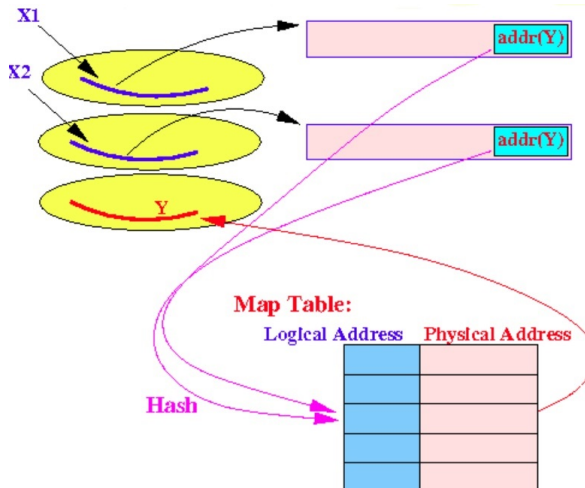
- When the tuple *Y* is deleted



- the physical addresses will reference an incorrect tuple


Techniques to handle tuple deletion

- Using logical addresses is easy
- Before deleting tuple Y that is referenced by tuples X1 and X2



Techniques to handle tuple deletion

- Using logical addresses is easy
- After deleting tuple Y



- Deleted tuple is identified by a NULL physical address in the Map table

Very important

- The logical address used by tuple Y must remain in the map table
- Furthermore:
 - The logical address used by tuple Y cannot be re-used

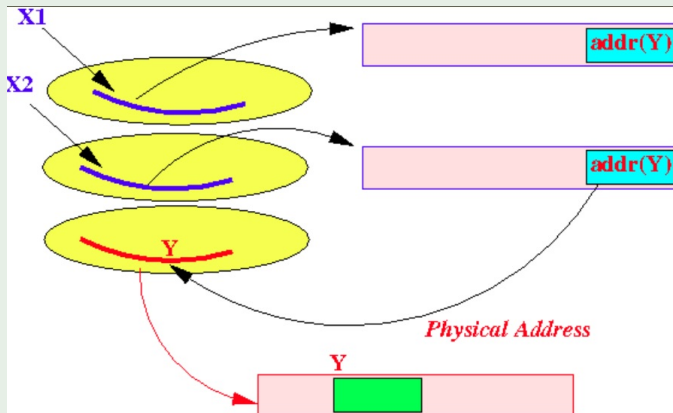
Techniques to handle tuple deletion

- Deleting a tuple using physical address: use a **tombstone** record
- **Tombstone record**
 - A tombstone record is a special-purpose tuple that is very small in size and serves as a marker to indicate that a tuple has been deleted.
- When a tuple is deleted, it is replaced by the **tombstone** record
 - Instead of physically removing the tuple from the database or shifting other records, you replace the target tuple with a tombstone record
 - This tombstone record essentially **marks** the tuple as deleted without removing it entirely.
- This **tombstone** is permanent, it must exist until the entire database is reconstructed
- *A map table or mapping structure is used to keep track of the relationships between logical or physical addresses and the actual data. When a tuple is deleted, the tombstone record can be represented in the map table, often as a null pointer or a specific value indicating deletion, in place of the physical address.*

Tombstones

Example


- Before deleting tuple Y



Tombstones


Example

- After deleting tuple Y



Tombstones

- When you insert a new tuple, you cannot use the space of a **tombstone** tuple (tombstone tuple must be preserved)
- Because: Existing tuple references to the deleted tuple will then references to the newly inserted tuple:



Update

- If new tuple is shorter than previous, easy
- If it is longer, need to shift tuples, create overflow blocks
- Note: We will never create a tombstone tuple in an update operation

Conclusion

- The storage manager is not entirely independent from the rest of the DBMS.
- A DBMS encodes and decodes the tuple's bytes into a set of attributes based on its schema.
- It is important to choose the right storage model for the target workload:
 - OLTP = Row Store
 - OLAP = Column Store

Database Storage: Next

- Problem#1: How the DBMS represents the database in files on disk
 - i.e., how to lay out data on disk.
- Problem#2: How the DBMS manages its memory and move data back-and-forth from disk.